


Choose the best fleet trajectory, not just the best next action.

ONE LIVE FLEET STATE · MANY CANDIDATE SEQUENCES · ONE FEASIBLE NEXT ACTION RETURNED



Aurelius adds a pre-execution planning layer above Kubernetes, Slurm, serving schedulers, autoscalers, and fleet-management systems. Before a decision is committed, it generates constraint-aware coupled-policy candidates, advances each against isolated simulated state under the configured constraints, rejects infeasible paths, and returns the highest-scoring feasible policy to the existing control plane. It reads scheduler and fleet metadata only, and never writes directly to the live fleet. In separately tractable spaces, the same bounded-search method matched exhaustive enumeration.

8.24× mean SLA-safe goodput per modeled infrastructure dollar versus a reconstructed production-class baseline	84.5% fewer GPU-hours under identical input request sets
	1,549,249 requests across three selected uncapped high-load windows
	3 of 3 windows with lower SLA violation rates

Simulated, load-dependent evidence from selected public-trace windows. Not a production result and not an estimate of performance against a particular fleet or internal scheduler.

The optimizer cannot choose a future it never represents.

The uncapped ablation reran the exact harness, windows, and load behind the 8.24× headline. The simulator, objective, cost model, constraints, and inputs remained fixed; only the planner's representable policy space and search procedure varied.

FIG 02 SEARCH-STRATEGY ABLATION · UNCAPPED REPLAY		FIXED HARNESS · POLICY SPACE + SEARCH		
SEARCH SPACE	COMPLETION	RELATIVE GP/\$	DELTA	MECHANISM
RECONSTRUCTED PRODUCTION-CLASS BASELINE	COMPLETES	1.00×		REACTIVE REFERENCE
GPU CLOCK ONLY	TIMES OUT*	0.69×*	-31.3%*	SINGLE SURFACE
FIXED 24-POLICY GRID	COMPLETES	3.62×	+261.9%	JOINT POSTURE
REGIME-SPECIFIC CANDIDATE SET	COMPLETES	3.62×	+261.9%	GENERATED, NO SEARCH
EXHAUSTIVE SEARCH OF FIXED GRID	COMPLETES	4.84×	+383.6%	FULL GRID ENUMERATION
BOUNDED COUPLED SEARCH	COMPLETES	4.85×	+385.4%	BOUNDED BEAM
HIERARCHICAL COUPLED-POLICY SEARCH (INDEPENDENT RERUN)	COMPLETES	8.29×	+728.6%	OFF-GRID COUPLING
SLA VIOLATION RATE 0.038 TO 0.046 AT THE BASELINE, 0.0073 TO 0.011 FOR THE GRID ARMS, 0.0022 TO 0.0026 FOR THE HIERARCHICAL SEARCH · LOWEST OF ANY ARM. * TIMES OUT AT THE STANDARD HARNESS BUDGET; AT AN EXTENDED BUDGET IT COMPLETES BELOW THE BASELINE.				
IDENTICAL INPUTS, WINDOWS, AND LOAD IN EVERY ARM · 1,549,249 REQUESTS				RERUN · 07.2026

The independent rerun reached 8.29× versus the canonical 8.24× benchmark, a 0.6% difference in the mean; individual market results were within 1.2%. **The report retains 8.24× as the canonical headline.**

A clock-only, single-surface search cannot be scored within the harness budget: the policy it selects degrades the simulated fleet until the replay exceeds it, in all three markets. Given far more compute it does complete, and still lands below the baseline (about 0.7×) with roughly ten times its SLA violation rate.

Setting the posture jointly is worth roughly 3.6× on its own: a fixed 24-policy grid and the regime-specific candidate set, scored without coupled search, select the same policy at this load.

Searching the grid harder is saturated: exhaustive enumeration of the fixed grid reaches 4.84×, and the bounded coupled search matches it at 4.85× while evaluating far fewer candidates.

The hierarchical coupled-policy search reaches 8.29× with the lowest SLA violation rates of any arm. The gain beyond the grid-based strategies comes from additional valid coupled-policy combinations that the fixed grid did not contain.

The largest gain came from expanding the policy combinations the planner could represent, not from searching the fixed grid more aggressively.

Uncapped search-strategy ablation, 07.2026: an independent rerun of the exact harness, windows, and load behind the headline benchmark, with production anchors reproducing the published values exactly. Ratios are unweighted means across the three markets. The frozen V1 audit separately measured zero search regret where exhaustive enumeration was tractable, at approximately 40 candidate evaluations per decision. Neither validates the simulator or establishes production savings.

Validating the benchmark result on operator data.

FIG 03 UNCAPPED REPLAY · SLA-SAFE GP/\$ VS THE RECONSTRUCTED PRODUCTION-CLASS BASELINE THREE MARKETS



1,549,249 total input requests · 84.5% fewer total GPU-hours · lower SLA violation rates in all three windows · both policies received the identical input request sets.

The comparison policy is a reconstructed production-class baseline: a deterministic public-replay policy assembled from established serving and cluster-control mechanisms, including continuous batching, deadline-aware ordering and admission, prefix-cache-aware routing, topology-aware placement, and backlog-driven capacity control. It is stronger than the naive SLA-aware reference, but it is not an operator's internal scheduler and does not reproduce a particular fleet's adaptive logic, internal signals, or cost model.


WHAT THIS ESTABLISHES

Under the selected uncapped high-load replay conditions, coupled predictive planning found feasible trajectories with materially better modeled economics and lower SLA violation rates than the reconstructed production-class baseline.

WHAT THIS DOES NOT ESTABLISH

Expected performance against a particular fleet, internal scheduler, workload distribution, or operator cost model. The result is simulated and load-dependent. Its magnitude is materially affected by the benchmark's batching, precision, service-time, and modeled cost assumptions.

Test one bounded historical window.



PROPOSED VALIDATION

Use one bounded historical window to estimate how much farther-ahead planning could have improved goodput per dollar under the fleet's actual workload and constraints.

Read-only metadata only. No payloads, production writes, or unrestricted infrastructure access.