

Optimizing the fleet's entire trajectory, not just its next decision.

A supervisory layer that forecasts the fleet forward, simulates many coupled futures under constraints, and returns one policy for the existing stack to execute.

Aurelius sits above the scheduler. It does not replace Kubernetes, Slurm, or a serving router, and it never touches payloads. Each control period it forks the current fleet state into many counterfactual trajectories, one per candidate coupled policy, advances each under SLA, capacity, power, and policy constraints, discards the ones that violate, and hands the highest-scoring feasible policy back to the stack to run.

+724% average SLA-safe goodput per dollar

MEAN OF +698% / +718% / +755% (PJM / ERCOT / CAISO), AN 8.24× GP/\$ RATIO VS A RECONSTRUCTED PRODUCTION-CLASS BASELINE · UNCAPPED REPLAY OF PUBLIC PRODUCTION TRACES (1.55M REQUESTS) · PARETO-SAFE AT ~84% FEWER GPU-HOURS

A number that size demands an explanation, and the report gives two. The **regime**: under heavy load a fixed-policy scheduler loses efficiency while a trajectory optimizer keeps adapting. The **mechanism**: on the same benchmark, with the simulator, objective, cost model, and gates held fixed across every arm, the result climbs from 3.62× the baseline on a fixed, predefined set of policies to the full headline only once the planner can form additional valid combinations of the same implemented controls, not only the predefined set. The increase came from what the architecture could represent, not from changing the simulator.

UNCAPPED HIGH-LOAD REPLAY · SIMULATED, EVIDENCE NOT A GUARANTEE · READ-ONLY, METADATA ONLY, NO PAYLOADS

00 Abstract

Conventional infrastructure stacks, including the unified control planes that consolidate a fleet into one view, optimize the next decision: the next placement, the next scale action, the next route, each chosen locally against the state visible now. Aurelius changes what gets optimized. It treats the fleet as a system with a trajectory, forecasts that trajectory forward, and optimizes the path rather than the step.

Concretely, it is not a forecaster and not a scheduler replacement. It is a supervisory layer that forks the current fleet state into many counterfactual trajectories, applies coupled policies across the controls the stack already exposes, advances each trajectory under hard constraints, kills the ones that violate, and returns one policy for the existing scheduler to execute.

In an uncapped replay of three selected public production-trace windows totaling **1.55 million requests** (the PJM, ERCOT, and CAISO expensive-power windows), Aurelius reaches an **8.24× mean gp/\$ ratio** against a reconstructed production-class baseline, where **gp/\$** is SLA-safe goodput, SLA-compliant served tokens, per infrastructure dollar. That is a **+724% average** gain (per market +698% / +718% / +755%), while holding SLA strictly better and using **~84% fewer GPU-hours**. It is a Pareto win, not goodput bought by spending more. The magnitude is load-dependent and requires operator-data validation.

A result that large is only worth reading if it is explained and defended, so this report does both. §07 gives the regime that produces the magnitude, a fixed policy losing efficiency under heavy load while a trajectory optimizer keeps adapting. §05 gives the mechanism, measured directly: on the exact uncapped benchmark, with the model, objective, cost model, and gates held fixed across every arm and only the planner's representable policies and search varied, the result climbs from a single-lever policy that cannot even complete the replay, to **3.62×** the baseline on a fixed, predefined set of policies, to the full **8.24×** headline once the planner can form additional valid combinations of the same implemented controls. The value was not in the model or the objective. It was in what the architecture could represent and search.

Scope, stated once and meant throughout: these are deterministic simulated replays on public production traces. The headline ratio is load-dependent and is quoted only with its scope (§07). This is directional evidence that requires live-telemetry calibration before any production claim, not a deployment. Throughout, Aurelius runs read-only and needs only scheduler metadata, never payloads or model weights.

01 The shift is in what you optimize

Most schedulers, and the control planes above them, optimize a current-state or local objective: availability, fairness, utilization, latency, queue state, immediate capacity, immediate price. Each is reasonable in isolation. But the economic outcome of a decision is set by constraints that arrive **after** it is made: power and capacity headroom, congestion, memory and topology pressure, demand, and the price of power, which on wholesale markets can move several-fold within a day.

Consolidating telemetry into one control plane is necessary, and it is not the missing piece. A unified present answers the question what is the best next decision given everything I can see now. It does not answer what is the best sequence of decisions given what is about to happen. Those are different questions, and the second is where the compounding waste lives, because once a job lands its costs are largely locked in and an observe-then-decide loop cannot price a constraint it has not yet seen.

Aurelius closes that gap by moving from an **observe, decide** loop to an **observe, forecast, simulate, decide** loop. It is a different control architecture, not another scheduler, and it changes the question from the next decision to the trajectory.

02 What Aurelius does each control period

Every control period, Aurelius forecasts the workload and conditions, forks the current fleet state into many counterfactual trajectories, one per candidate coupled policy, advances each trajectory one control period at a time against the forecasted constraints, rejects any that fail a hard gate, scores the survivors by risk-adjusted SLA-safe goodput per dollar, and returns the highest-scoring feasible policy for the existing scheduler to execute. When forecast confidence is low, it falls back deterministically to the stack's default.

The word that carries the architecture is **coupled**. Aurelius does not tune one knob. Workload timing, placement, routing, capacity, batching, and the precision and clock policy are scored together as one bundle, because those decisions interact, and, as the finding shows, the interaction is exactly where the economic value lives.




FIG.01 · ONE PLANNING CYCLE: MANY COUNTERFACTUAL TRAJECTORIES FORKED, ADVANCED UNDER CONSTRAINTS, ONE RETURNED · SCHEMATIC, NOT TO SCALE

03 The predictive world model

Aurelius maintains a persistent world model of the cluster: a forward model of the fleet, not a snapshot of free GPUs. It carries replica warm and cold state, server and rack capacity, placement and locality, queue dynamics, in-flight migrations, and the operator-cost basis, and it advances that state one control period at a time so a candidate decision is scored against the state it will actually meet.

Candidate policies are evaluated on read-only simulated future states, so scoring never mutates the live cluster. Every input is tagged by provenance, measured, trace-derived, benchmark-calibrated, or modeled, so no result silently treats a modeled value as a measured one.

04 Which futures get evaluated is a first-class step

Because the controls interact, Aurelius evaluates coupled bundles rather than tuning one knob, and which bundles it evaluates is itself a first-class step, not an afterthought. A forecast-guided generator

proposes a focused set of high-value coupled bundles implied by the operating regime, always including the known-strong candidates so a high-value option is never silently dropped. The set is then searched in a way that captures cross-surface coupling, widening to more surfaces only when a decision is close and stopping early when the choice is clear.

Where exact enumeration is tractable, the bounded search is scored against it, so any regret is **measured** rather than assumed. That discipline is what surfaced the finding in §05, and it is what lets a several-fold headline be defended rather than asserted.

ACTIVE DECISION LEVERS

SCORED AS COUPLED BUNDLES; EACH CHANGES THE PROJECTED ECONOMIC OUTCOME

Capacity policy · Capacity multiplier · Ordering · Admission · Batching · Routing · Prewarm · Placement · Migration · Precision · Clock / DVFS

05 The limit was representation, not search effort

The thesis of this report is a claim about where orchestration value actually sits, and it is not in a better next decision. Most production systems optimize one decision at a time and re-plan after it executes. Aurelius evaluates the consequences of a sequence of coupled decisions before committing the first one: it searches for the optimal **path** of coupled decisions, forecast many control periods ahead under constraints, and commits only the first step of the best path it finds. That search, over forecasted trajectories rather than the state visible now, is where the value current orchestration leaves on the table is recovered.

The claim is testable, and the test is controlled. On the exact uncapped benchmark that produces the headline in §06, the forecast, simulator, objective, cost model, gates, replay inputs, and baseline remained fixed. Only two things varied: the planner's representable policy space and the procedure used to search it. Every arm sees the same world and is judged the same way, so the headline gap can be pinned to representation and search, not to a weaker baseline or a more generous model.

TABLE 1 · SAME UNCAPPED HARNESS, VARYING ONLY REPRESENTABLE POLICIES AND SEARCH · SLA-SAFE GOODPUT/\$ RELATIVE TO THE RECONSTRUCTED BASELINE (MEAN ACROSS PJM, ERCOT, CAISO)

SEARCH STRATEGY	GP/\$ VS BASELINE	SLA VS BASELINE	NOTE
Reconstructed production-class baseline	1.00×	–	reference anchor
Single lever (clock only)	does not complete	–	replay cost blows up under load
Fixed set of 24 predefined policies	3.62×	~4× lower	coupling, no search
Search that fixed set (bounded)	4.85×	~4× lower	matches full enumeration of the set
Expanded coupled-policy space, same controls	8.24×	lowest, ~18× lower	the +724% headline arm

Every arm is scored against the same byte-identical baseline (130,538.91 / 130,178.26 / 130,000.36 gp/\$ across the three markets); the coupled-search arm is the §06 headline configuration, quoted at the headline value (8.24×). Simulated; every completing arm also lowers the SLA-violation rate, the coupled arm the most (to ~0.002).

A word on the two settings compared, because the distinction is the whole point. The **fixed set of 24 predefined policies** is a discrete set enumerated in advance over four implemented controls: batching, capacity, precision, and clock. Constructing **additional valid combinations of those same implemented controls**, the wired levers in §04, introduces no new or hypothetical control; these are

configurations the 24 predefined policies did not include. Every policy, inside the set or outside it, runs through the identical simulator, cost model, and safety gate, so nothing outside the set is unpriced or unconstrained.

Read top to bottom, the ladder says three things:

- **A coherent joint policy already produces a large gain.** The fixed set of 24 predefined policies evaluates batching, capacity, precision, and clock jointly rather than independently. That alone reaches 3.62× the reconstructed baseline and reduces SLA violations by roughly fourfold.
- **More search over the same fixed set adds little.** Bounded search over those 24 policies reaches 4.85×, essentially matching exhaustive evaluation of the same set at 4.84×. The planner is therefore already finding the best policy available within that predefined candidate set.
- **The remaining gain comes from expanding the representable policy space.** The increase from 4.85× to 8.24× appears when Aurelius can construct additional valid combinations from the same implemented controls rather than being restricted to the original 24 policies. The limiting factor was not search effort. It was which coupled policies the planner was able to represent and evaluate.

The result is controlled in two ways. The increase was not created by changing the simulator, objective, cost model, or constraints, which remained fixed across every arm; it came only from expanding which coupled policies the planner could represent and evaluate. And where a window is small enough to enumerate the entire space, that same bounded search has matched the exhaustive optimum exactly, so the value is captured rather than stumbled onto. Whether the simulator's absolute magnitude transfers to a real fleet is a separate question, and the purpose of the historical replay proposed in §11.

One exclusion, stated plainly. An aggressive low-precision (int4) mode scores higher still, but its quality cost is not yet modeled, so it is excluded from every headline and reported only as a labeled diagnostic ceiling. The headline-safe result uses the precision band whose quality can be defended.

06 The headline result: uncapped replay at production-scale volume

The headline replays the full selected windows with the per-period request cap removed, so each runs at its natural, production-scale request volume: PJM ~577k, ERCOT ~443k, CAISO ~530k, about 1.55M in total. Both policies run the identical load. Because no operator's internal scheduler is publicly available, I reconstructed a production-class baseline from established components used in modern serving stacks: continuous (iteration-level) batching^{[1][2]}, SLA-aware ordering and admission^[3], prefix-cache-aware routing^[4], topology-aware placement^[5], backlog-driven autoscaling^[6], and warm-pool headroom. It is not any particular operator's proprietary scheduler; it is the best reproducible approximation of a capable production stack using public mechanisms and data. Defined once here, it is called the **production-class baseline** from now on, and the comparison isolates the effect of how Aurelius combines and searches the same implemented controls.

At that volume the naive SLA-aware policy does not complete: its no-batching, no-admission replay cost grows super-linearly and times out, so it is a reference only. The production-class baseline and Aurelius both complete, which makes them the primary, production-comparable comparison: the policy built from real serving-stack components is the bar Aurelius has to beat.

MARKET	AURELIUS GP/\$	PRODUCTION GP/\$	Δ %	SLA (AUR.)	SLA (PROD.)	REQUESTS
PJM · expensive	1,041,842	130,539	+698.1%	0.0023	0.0382	576,912
ERCOT · expensive	1,064,602	130,178	+717.8%	0.0013	0.0447	442,716
CAISO · expensive	1,111,915	130,000	+755.3%	0.0018	0.0458	529,621
Average / total	–	–	+723.7%	–	–	1,549,249

gp/\$ and SLA-violation rates read directly from run output. Δ % is per market (Aurelius – production) $\div$ production; the +724% headline is the equal-weighted average of the three (723.7%, shown unrounded). SLA violations fall from ~3.8–4.6% to ~0.1–0.2%.

TABLE 3 · GPU-HOURS AT THE SAME UNCAPPED LOAD

MARKET	AURELIUS GPU-H	PRODUCTION GPU-H	Δ GPU-HOURS
PJM · expensive	62.2	399.9	–84.4%
ERCOT · expensive	51.0	323.7	–84.3%
CAISO · expensive	60.0	391.5	–84.7%
Total	173.1	1,115.1	–84.5%

Same served requests both arms; the reduction is (production – Aurelius) $\div$ production. Aurelius earns more SLA-safe goodput on far fewer GPU-hours because the coupled bundle packs more work per GPU-hour while holding SLA, which is why the win is Pareto and not paid for in spend.

07 How to read a several-fold number

A sevenfold gap is unusual, so the report is explicit about the two effects that compound to produce it, and about where its edges are.

Two effects that compound. Under the selected heavy load, the fixed-policy baseline loses efficiency as the replay becomes saturated, its goodput per dollar falling toward a similar floor in each market because a fixed policy hits the same saturation under the same load, while Aurelius continues adapting its coupled policy. The ablation in §05 then shows that the largest incremental increase appears when the planner is no longer limited to the predefined policy set. Both effects are real, and neither alone is the whole number.

Where its edges are. The multiple is load-dependent. It moves with the per-period request volume and with how stressed the fleet is, so it is quoted only with its uncapped, full-window scope, never as a fleet constant. What does not depend on the load is the attribution in §05: the value comes from representing and searching coupled futures, and the direction holds across every load tested.

Read the several-fold number as two things at once: a stress signal, showing how far a fixed policy falls behind under load, and a mechanism, the coupled search that reaches policies a fixed, predefined set cannot express. It is not a fleet constant, and this report never quotes it as one.

08 What is actually new here

The individual pieces are not new, and this report does not claim they are. Forecasting, model-predictive control, digital twins of infrastructure, SLO-aware scheduling, and unified control planes all exist and are well understood. Aurelius's claim is narrower, and more durable for it.

What Aurelius optimizes is the fleet's **forecasted trajectory**, evaluated as coupled futures rather than independent next-decisions, and the binding constraint on the result is not the model or the objective but **which coupled futures the planner can represent and evaluate**. The ablation in §05 is the evidence: on an identical stack, varying only which policies the planner could represent and how it searched them moved the result from a policy that cannot complete the replay, to 3.62× the baseline on a fixed, predefined set of policies, to 8.24× once the planner could form additional valid combinations of the same implemented controls. That is a claim about where the value lives in fleet optimization, and it is testable on any fleet that keeps historical scheduler metadata.

09 Safety and the deployment path

Aurelius is **read-only and metadata-only**. It reads only the metadata a scheduler already exposes, job timing, resource requests, constraints, and never payloads, model weights, or model outputs. Candidate scoring runs on simulated futures, so it writes nothing to the live cluster. Unsafe candidates are rejected at the constraint gate before any recommendation exists, and when confidence is low the controller falls back deterministically to the stack's default.

The path from telemetry to any production change is staged and reversible: historical replay of past decisions first, then recommendation-only shadow mode alongside the live scheduler, then optional actuation on workloads explicitly in scope. Aurelius does not directly write fleet state. It returns a constrained recommendation through the existing control plane, which remains authoritative and owns execution.

10 Limitations

- Every number here is deterministic simulated replay of public traces: directional evidence of achievable savings, not a guarantee for any specific fleet, and not a production deployment.
 - The +724% headline is load-dependent. Part of the gap is a fixed-policy scheduler degrading under heavy uncapped load while Aurelius adapts, so it is reported only with its uncapped, full-window scope and never as a fleet constant.
 - Magnitudes are simulator-inferred and depend on the serving model's precision and batching bands. The most robust findings are the **direction** and the **attribution to search**, not any single percentage.
 - Aggressive low-precision is excluded from every headline until its quality cost is modeled.
 - Simulator fidelity must still be tested against real operator telemetry, which is the only way to isolate model error, and the largest gains appear during expensive-power windows.
 - Some controls are modeled but not yet active production levers, and several constraints are represented but not yet forecast inputs.
-

11 Validating the finding on a real fleet

The next step is not a bigger simulation, it is calibration against reality. The baseline here is the strongest public approximation I could build; the next honest test is not another synthetic baseline, it is replay against an operator's own historical decisions. Given read-only historical scheduler metadata, Aurelius would replay that fleet's recorded behavior, simulate the coupled counterfactuals, and produce an audited estimate of the SLA-safe goodput per dollar left on the table, with no production changes and no payload access.

The test has a built-in honesty gate. Before scoring any counterfactual improvement, Aurelius first has to reproduce the operator's recorded baseline behavior within agreed tolerances. If it cannot

reconstruct that behavior and improve a pre-agreed metric on held-out telemetry, the test ends. If it can, the next step is recommendation-only shadow mode. That is the fastest, lowest-risk way to find out whether the result generalizes, and the quickest way to prove it wrong.

Reproducibility notes & references

BENCHMARK CONFIGURATION

- **Primary metric.** SLA-safe goodput per infrastructure dollar. The numerator is SLA-compliant served tokens, tokens served within their latency or deadline target; work served outside that target is excluded, so goodput bought by breaking SLA does not count. The denominator, modeled infrastructure cost, is the sum of three terms: GPU cost (active GPU-hours, including warm-held replicas, at a fixed hardware cost basis), energy cost (GPU energy at the prevailing wholesale electricity price), and network cost (including migration). The GPU cost basis is hardware only, so energy is counted once, at the market price. This is a modeled benchmark cost, not full operator TCO.
- **Windows and load.** The expensive-power windows of PJM, ERCOT, and CAISO, replayed uncapped at full volume (~577k / ~443k / ~530k requests, ~1.55M total), three control-period decisions per window.
- **Configuration.** Deterministic replay under a fixed seed; every policy is scored on the same windows under the same simulator, cost model, and constraint gates. Forecasts use only history up to each decision, and per-request execution time is replayed rather than predicted, so no policy sees the future. The baseline's goodput per dollar reproduces byte-identically across runs (130,538.91 / 130,178.26 / 130,000.36 gp/\$). At each decision the hierarchical search evaluates roughly 75 coupled policies out of a representable space of about 4.7M.

The reconstructed production-class baseline (§06) is assembled from established production serving and cluster-scheduling techniques ([1]–[6]), cited for the realism of the baseline mechanisms, not for any Aurelius result. The public workload and electricity-price data are cited at [7] and [8].

[1] Yu et al. “Orca: A Distributed Serving System for Transformer-Based Generative Models.” OSDI 2022. Iteration-level (continuous) batching.

[2] Kwon et al. “Efficient Memory Management for Large Language Model Serving with PagedAttention” (vLLM). SOSP 2023. Continuous batching and KV-cache management.

[3] Gujarati et al. “Serving DNNs like Clockwork: Performance Predictability from the Bottom Up.” OSDI 2020. SLO-aware scheduling and admission.

[4] Zheng et al. “SGLang: Efficient Execution of Structured Language Model Programs” (RadixAttention). 2024. Prefix / KV-cache-aware reuse.

[5] SchedMD. “Slurm Workload Manager, Topology Guide.” Topology-aware placement in production cluster schedulers.

[6] Rzadca et al. “Autopilot: Workload Autoscaling at Google.” EuroSys 2020. Utilization-headroom autoscaling at production scale.

[7] Workload, public serving traces: Wang et al. “BurstGPT: A Real-world Workload Dataset to Optimize LLM Serving Systems” (2024), for arrivals and request characteristics; and the Mooncake trace, for prefix / KV-cache reuse.

[8] Electricity, public wholesale prices: CAISO OASIS, PJM Data Miner, and ERCOT, day-ahead and real-time locational marginal prices for the selected PJM, ERCOT, and CAISO windows.